


Deep Learning for EEG Seizure Detection

Project Report



Abstract

- **4 models**: threshold, random forest, CNN, LSTM
- **Patient-level** and **seizure-level** cross-validation (5-fold)
- CNN: 86.3% acc (patient-level); LSTM: **97.2% acc, 0.985 F1** (seizure-level)
- Window-level CV produces degenerate results

Contents

1. Introduction & motivation
2. Dataset & analysis
3. Preprocessing & data flow
4. Methodology: 4 pipelines
5. Cross-validation strategies
6. LSTM design: pooling & context
7. Hyperparameters & experiment grid
8. Results & conclusion

Introduction

Epileptic seizures = abnormal brain electrical activity detectable via EEG.

Goal: automatically detect seizures in EEG recordings.



Approach: compare progressively richer models:

1. Threshold baseline (hand-crafted rules)
2. Random forest (hand-crafted features)
3. CNN (learned features, end-to-end)
4. LSTM (temporal episode modelling)

Dataset: CHB-MIT


- **Source:** CHB-MIT scalp EEG, 24 pediatric patients
- **Windows:** 1s segments → **571,905** total
- **Seizure:** **86,672** windows (15.2% — imbalanced dataset)
- **Shape:** each window = **21 channels × 128 timesteps**

Dataset: Feature Trends




- **Std** and **range** separate seizure from non-seizure
- **Mean** ≈ 0 for both $\rightarrow$ risky feature alone

Preprocessing & Data Flow




- No per-window normalization (`normalize=False`)
- **Class weighting** helps models focus on the seizure minority

Overview: 4 Pipelines




Pipeline 1: Threshold Classifier




- Computes **std, range, min, max** per window
- Flags seizure if stat exceeds fold-specific threshold
- Thresholds = midpoint of seizure/non-seizure means

Pipeline 2: Random Forest



- **8 features:** mean, std, min, max, range, ptp, std/range ratio, range+std
- 200 trees, `class_weight='balanced'`
- Patient-level 5-fold CV


Pipeline 3: CNN



- **Input:** raw window [21, 128], no normalization
- 3× Conv1d → BN → ReLU → MaxPool
- FC: 2048→256→Dropout(0.3)→2
- Adam (lr=1e-3), class-weighted CE, 2 epochs

Pipeline 4: LSTM

*Adapted from course-provided EpilepsyLSTM
(ChakrabartiChannelFusion)*



- **Per-timestep classification** (original: last hidden state only)
- **Conv1d projection** option + packed sequences for variable-length episodes
- Per-window **pooling** (std/mean/mean_std) reduces [21,

Cross-Validation Strategies

Seizure-level

Whole episode → one fold

± 20 context windows in val

Window-level

Window round-robin

± 4 seizure neighbors
dropped

Seizure-Level k-fold: Context

- Without context: val = **100% seizure** → no specificity measurable
- With `context=20`: ± 20 non-seizure windows around each episode
- Val gets realistic $\sim 40\text{--}60\%$ non-seizure mix
- Train gets all remaining windows

LSTM: Pooling Options

EEG oscillates around 0 → **mean** collapses amplitude. **Std** preserves it.

	Pool	Dim	Speed	Preserves
	std (default)	21	Fast	Channel variance
	mean	21	Fast	Near-zero — risky
	mean_std	42	Fast	Center + spread
	conv_proj	32	Medium	Learned features
	none	2688	Very slow	Full waveform

LSTM: Architecture

```
Input: [batch, seq_len, n_features]
      |
      v
LSTM(input=n_features, hidden=128, layers=2, dropout=0.3)
      |
      v
Linear(128 → 2) per-timestep
      |
      v
Output: [batch, seq_len, 2]
```

- Packed sequences for variable-length episodes
- Class-weighted CE, `ignore_index=-1` for padding

LSTM: conv_proj Mode

Conv1d projection option replaces hand-crafted pooling:

```
[seq_len, 21, 128]
|
Conv1d(21 → 32, k=5, pad=2) → BN → ReLU
|
AdaptiveAvgPool1d(1)
|
[seq_len, 32] → LSTM(input_size=32)
```

- Learns waveform features end-to-end
- Faster than **none** (2688-d), more expressive than **std** (21-d)

Hyperparameters


Parameter	Threshold	RF	CNN	LSTM
Trees/layers	—	200	[32,64,128]	2 LSTM
Hidden size	—	—	—	128
Pooling	—	—	—	std / conv_proj
Class weight	—	balanced	balanced	balanced
LR	—	—	1e-3	1e-3
Batch size	—	—	32	32
Epochs	—	—	2	2
Dropout	—	—	0.3	0.3
CV level	patient	patient	patient	seizure (ctx=20)

Experiment Grid

Script	Runs	~Time
exp1_cnn_levels	CNN × 3 levels	~1.5h
exp2_lstm_levels_pool	LSTM × 3 levels + 4 pools	~5-6h
exp3_lstm_nsl	LSTM × 4 NS lengths	~3.5h


All write to `out/benchmarks/` with descriptive suffixes.

Results: Accuracy per Fold (Patient-Level)



High cross-patient variance across all models;
CNN ranges 79–93%.

Results: Main Comparison (5-fold)




Model	CV	Acc	F1
Threshold	patient	75.5%	—
RF	patient	69.0%	—
CNN	patient	86.3%	0.554
CNN	window	97.7%	0.000
CNN	seizure	80.1%	0.888
LSTM-std	seizure	97.2%	0.985
LSTM-std	patient	99.5%	0.983

Results: Window-Level CV Degeneracy

CNN and LSTM under window-level CV: **97-100% accuracy but 0% F1**

- Round-robin assignment creates unreliable per-fold class ratios
- Models learn to predict "non-seizure" for everything → trivially high accuracy
- **Window-level CV is unsuitable** for this dataset

Results: CNN Split-Level Comparison



CV Level	Avg Acc	Avg F1
patient	86.3%	0.554
window	97.7%	0.000
seizure	80.1%	0.888


“

Patient-level:
meaningful but
variable.

Window-level:
degenerate.

Seizure-level: high


Results: LSTM Pooling (Seizure-Level)



Pooling	Avg Acc	Avg F1
std (21-d)	97.2%	0.985
mean (21-d)	97.8%	0.988
conv_proj (32-d)	93.3%	0.963
mean_std (42-d)	61.2%	0.601

“ mean_std unstable
(2/5 folds fail)

Results: Non-Seizure Length Sweep



NSL	Avg Acc	Avg F1
5	98.8%	0.994
10	96.1%	0.979
20	5.5%	0.000
50	5.5%	0.000

“

NSL $\geq$ 20 $\rightarrow$
training collapse

”

Key Observations

- **LSTM (seizure-level)** dominates: 97.2% acc, 0.985 F1
- **CNN (patient-level)**: 86.3% acc but F1 only 0.554 — low recall on some patients
- **Window-level CV** produces degenerate results for all models
- **Threshold (75.5%)** outperforms RF (69.0%)
- **Mean & std pooling** both work well; mean_std is unstable
- **NSL ≥ 20** causes LSTM training collapse

Conclusion

- LSTM (seizure-level) achieves best results: **97.2% acc, 0.985 F1**
- CNN (patient-level): 86.3% acc, 0.554 F1 — high variance across patients
- Window-level CV is degenerate — poor class ratios per fold
- Mean pooling slightly outperforms std for LSTM
- **NSL ≤ 10** is essential; NSL ≥ 20 causes collapse
- **Future:** explore multichannel spatial info, longer temporal context

Thank You

Any Questions?